



Ministry of Housing,
Communities &
Local Government

Refining the planning data specification

Data Design Team

September 2026

Table of contents

Introduction.....	2
What we're working towards and why	2
Understanding the specification	3
Clearer guidance and practical examples	5
Using it now and helping it improve	6
Credits	7

Introduction

The more people pick up, review and use the planning application data specification, the more we learn and the better it becomes. Recent work with planning authorities and software suppliers has helped us clarify field definitions, explain important distinctions and add practical examples, making the specification easier to use.

This report explains the approach and how to use the material now available. Some feedback led to changes to the specification. In other cases, the specification already met the need, but we needed to explain it better or make the supporting material easier to find.

What we're working towards and why

The 300+ planning authorities in England hold application information in many different systems. Making it available in a consistent structure helps people to compare, connect and reuse it without having to interpret each authority's records differently.

That matters for more than finding an individual application. Taken together, applications show what is happening across the country: where development is proposed, how applications progress and what decisions are made. This includes smaller proposals, such as extensions and loft conversions, as well as major developments.

The specification brings consistency to information recorded through the planning process. It does not yet cover everything that may form part of the statutory planning register, and we will continue to develop it as people bring forward further needs. Those [user needs](#) help us decide what information to include and why.

The specification, as it currently stands, connects records for applications, sites, decisions, documents, conditions and progress through the process. It is a foundation, not a claim that these are the only useful facts. We will add more information about what an application proposes, such as housing numbers or floorspace, to the specification as needs are established.

Adding more fields is not the only way to make the data more useful. It can also be combined with other planning data, including datasets available through [Planning Data](#). For example, comparing a site boundary with conservation areas or land allocated for development helps people understand the application's context without copying that information into every application record.

Making this possible may require authorities and suppliers to adapt their systems and how they record information. The specification gives that work a common direction: data that can support public access, reporting and analysis without being prepared separately for each purpose. We use feedback about existing systems to understand the changes needed, while keeping the focus on what users need from the data.

Understanding the specification


Connected records, not a single list of fields

Applications can propose development or relate to an earlier permission, for example to discharge a condition. The specification connects application records to the land they concern through a separate **site** record.

Each application has its own reference. Each site has its own reference.

The site reference identifies a separate site record, not an application or an individual property. Applications can be linked to the same site using the reference of the shared site record (see diagram below). A site record may include one, several or no UPRNs.

References are useful even when they are not prominently displayed on a public webpage. They are the glue that holds the data together. The specification describes data that services can use, not how every service must present it.



References connect related records without repeating the same information in each one.

[View a larger version of the references diagram](#)

If you want to see how this fits together, start with the [complete householder example](#). It brings the records together in one downloadable JSON file, with an introduction explaining their relationships. You can follow the references from the application to its site, and from the documents and decision notice back to the application.

Recording information is different from publishing it

The [datasets](#) describe information an authority is expected to hold through the planning process.

The [national public view](#) is part of the specification. It defines which records and fields are to be published as open data; it is not a separate platform or service.

This specification is not a design for a national planning-register website: that would be a service built using the data.

Nor does it mean publishing everything an authority holds.


The householder example includes a withheld document to illustrate this distinction. Each document held by the planning authority has its own publication assessment.

Under these rules, only document records marked **publish** are included in the open data. Records marked **withhold** or **not-assessed** are excluded.

This is a deliberate safeguard: the authority must explicitly assess a document as suitable for publication before its record is included in the published open data. Documents are not open by default. A document awaiting assessment stays outside the national public view.

Recorded does not mean published

Document records in the national public view



Only the specified records and fields are included in the open data. Documents awaiting assessment are not included.

Clear expectations about what to provide

We have added [requirement levels](#) to distinguish information that **MUST** be provided when applicable, information that **SHOULD** normally be provided and information that **MAY** be provided.

This avoids treating every field as either compulsory in every circumstance or entirely optional. For example:

- A site reference is essential for connecting records, so it is **MUST**.
- A meaningful site name helps people recognise the site, so it is **SHOULD**. Record one where it exists and is useful, but do not invent one merely to fill a field.
- Notes are optional, or **MAY**. Include them if they add useful context.

Requirement levels are set separately for each field in the datasets and the national public view. Holding information in an authority's system and publishing it as open data serve different

purposes, so the expectations can differ. Equally, information needed in the authority's records is not automatically required in the open data.

Clearer guidance and practical examples

The [dataset pages](#) now bring together field details, requirement levels, guidance and examples. Examples can be read as tables or JSON. The sections below cover areas people found confusing and show how we have used clearer guidance and practical examples to explain them.

Identifying sites and related applications

The [site dataset](#) records an address or description to help people identify the site and a boundary showing the exact land it covers.

The `address-text` field can contain a postal address or describe unaddressed land, such as "Land north of Station Road". We removed the separate description field because it duplicated the purpose of address text. We've also added property and street identifiers, UPRNs and USRNs. The [address discussion](#) provides more background on why we use this approach.

The `linked-applications` field identifies applications with a direct procedural relationship, rather than every application in a site's planning history. We have added [guidance with examples](#), such as linking a technical details consent application to its earlier permission in principle application, to show how to use it.

A submitted application can also cover more than one application type, as existing forms on GOV.UK already allow. For example, this [combined full and listed building consent example](#) shows how to record multiple types in the `application-types` field.

Following progress without losing important dates

The planning permission timeline dataset records milestones that help people understand how an application progresses, including validation and consultation. It is not a request for every phone call, email or internal action. The [event codelist](#) defines the events we want to be recorded, with a specified selection for the national public view. The [community discussion about the timeline and its events](#) explains the needs behind the original proposal; the codelist contains the current definitions. The [guidance and example](#) show how to use them.

The received date is a basic fact about the application, so it sits on the application record where anyone looking at that record can find when the authority first received it. We do not duplicate it in the timeline: [recording each date in one place](#) avoids conflicting values. Validation is a separate part of the process. Timeline events record when the application was treated as valid or invalid and when the validation check was completed. Together, these details tell the application's story.

These facts then help amongst other things to distinguish time spent reaching validation from time spent determining the application. They also help explain interruptions and delays, as set out in the [rationale for recording validation information](#).

Recording decisions and recommendations

Decisions and formal officer recommendations now share the values grant and refuse, via the [decision codelist](#). A [decision record](#) explains this choice.

Because conditions are recorded separately, we do not need a separate "grant with conditions" outcome. This can be inferred from the conditions attached to the decision.

We've given a withdrawal its own date field because this covers both that an application was withdrawn and when.

Identifying who made the decision

Knowing which authority issued a decision does not tell us whether an officer or a planning committee made it. That's what the `decision-maker` field records. It is the category or level of decision maker, not an individual officer's name. The organisation field in the decision notice dataset identifies the organisation issuing the notice. The [decision-notice guidance and examples](#) show how to record officer, committee and Planning Inspector decisions.

Explaining conditions and their reasons

For conditions, the requirement and the reason for imposing it remain separate. This makes both usable without extracting one from a combined passage. The condition examples illustrate that separation.

We have also clarified the different organisation roles and removed a duplicate organisation field. The [decision-condition example](#) distinguishes the body requesting a condition from the authority issuing the decision.

The `discharged-by` field identifies the later decision notice that discharged the condition, not an organisation.

Making documents easier to use

The document guidance and examples explain readable names, publication status and replacement. For example, a newer document points to the earlier (document) record it replaces.

We have also clarified the two URL fields. `document-url` should always link directly to a file, such as a decision notice PDF. `documentation-url` links to a webpage where someone can find information about the record and access related documents.

Using defined values consistently

Some fields use a codelist: a defined set of values with agreed meanings. Use these values rather than local system labels so data from different authorities can be understood consistently. The [codelists](#) and [application guidance](#) explain which values to use, including development scale and planning authority identifiers.

Using it now and helping it improve

Use the [specification viewer](#) to explore the datasets, field definitions and guidance. If you want to follow the connections between records, look at the [worked householder example](#).

The [examples library](#) also includes submission examples. Definitions, guidance and decision records provide progressively more detail without needing to be read together from beginning to end.

You can use the specification now and help shape its further development. Specific evidence helps us improve it while retaining choices that already meet established needs. We are seeking examples of [Section 106 agreements published across multiple documents](#) and [cross-boundary applications](#).

We will keep asking questions in community sessions and on [GitHub](#) as other topics arise. We want to draw on your wide-ranging expertise.

When raising an issue, identify the topic, the field or the rule, explain the difficulty and include an example where possible. This helps distinguish an unclear explanation, a genuine specification gap and work needed to implement an existing requirement.

Credits

Thank you to the planning authorities and software suppliers who contributed mapping returns, examples and discussions through the Development Management Software Fund work, and to the wider planning data community for helping improve the specification and its guidance.